AGI Strategy Reading List

Luke M. Perez

last updated: October 16, 2026

Unit 1: The Race to AGI

What's at Stake

Required Readings

- Selsam, Daniel. 2026. "Daniel Selsam - Personal Statement on AI Risk." Google Docs. September 14. https://docs.google.com/document/d/1NRSfwBrVGWkBE24zRKOeP_Ksmx8ZRLbIwsF1k8pjDfA/edit?

Optional Readings

- Karnofsky, Holden. 2021. "The Most Important Century (in a Nutshell)." Cold Takes. September 23. https://www.cold-takes.com/the-most-important-century-in-a-nutshell/.
- Amodei, Dario. 2024. "Machines of Loving Grace." October. https://darioamodei.com/essay/machines-of-loving-grace.
- MacAskill, William. 2025. "Introducing Better Futures." Forethought. August 3. https://www.forethought.org/research/introducing-better-futures.
- Altman, Sam. 2025. "The Gentle Singularity." Sam Altman. June 10. https://blog.samaltman.com/the-gentle-singularity.
- Gates, Bill. 2026. "The Choices We Make about AI Now Are Critical | Bill Gates." gatesnotes.com. August 26. https://www.gatesnotes.com/a-turbulent-ai-era-and-critical-choices-to-make.

What Future do you want?

Optional Readings

- Carlsmith, Joe. 2022. "Actually Possible: Thoughts on Utopia." Joe Carlsmith. October 12. https://joecarlsmith.com/2021/01/18/actually-possible-thoughts-on-utopia.
- Binder, Damon. 2026. "Post-AGI, We Are All Jobless Aristocrats." Defenses in Depth. September 11. https://defensesindepth.bio/post-agi-we-are-all-jobless-aristocrats/.

- Yudkowsky, Eliezer. 2009. "The Fun Theory Sequence." LessWrong. January 25. https://www.lesswrong.com/posts/K4aGvLnHvYgX9pZHS/the-fun-theory-sequence.
- Yudkowsky, Eliezer. 2021. "I'm from a Parallel Earth with Much Higher Coordination: AMA." LessWrong. April 5. https://www.lesswrong.com/posts/gvA4j8pGYG4xtaTkw/i-m-from-a-parallel-earth-with-much-higher-coordination-ama.

Competition and Coordination

Required Readings

- Karnofsky, Holden. 2022. "Racing Through a Minefield: The AI Deployment Problem." Cold Takes. December 22. https://www.cold-takes.com/racing-through-a-minefield-the-ai-deployment-problem/.
- Wildeford, Peter. 2026. "The Scramble: Getting in Position to Pace the Frontier." August 20. https://blog.peterwildeford.com/p/the-scramble-getting-in-position.

Recommended Readings

- Toner, Helen. 2025. "In Search of a Dynamist Vision for Safe Superhuman AI." Substack newsletter. Rising Tide. May 12. https://helentoner.substack.com/p/dynamism-vs-stasis.
- Samuel, Sigal. 2024. "It's Practically Impossible to Run a Big AI Company Ethically." Vox. August 2. https://www.vox.com/future-perfect/364384/its-practically-impossible-to-run-a-big-ai-company-ethically.
- Rehman, Iskander, Karl P. Mueller, and Michael J. Mazarr. 2025. "Seeking Stability in the Competition for AI Advantage." https://www.rand.org/pubs/commentary/2025/03/seeking-stability-in-the-competition-for-ai-advantage.html.
- "The Intelligence Age." 2024. September 23. https://ia.samaltman.com/?utm_source=bluedot-impact.
- Ó hÉigeartaigh, Seán. 2025. "The Most Dangerous Fiction: The Rhetoric and Reality of the AI Race." SSRN Scholarly Paper. Rochester, NY. May 25. doi:10.2139/ssrn.5278644.

Responding to the Hugging Face Incident

Required Readings

- Patel, Dwarkesh. 2026. "The Rise and Fall of Agent Civilizations." August 29. https://www.dwarkesh.com/p/openai-huggingface.

Optional Readings

- Patel, Dwarkesh. 2026. "Ajeya Cotra – Inside the OpenAI Agent Swarm That Hacked Hugging Face." September 1. https://www.dwarkesh.com/p/ajeya-cotra.
- Von Arx, Sydney, Cormac Slade Byrd, Spencer Kitts, and Thomas Larsen. 2026. "Discovery of a New OpenAI Agent Message Board." September 4. https://collusion.

wiki/.
- Arx, Sydney Von, Thomas Larsen. 2026. "OpenAI Agents Carried Out an Undisclosed Cyber-Attack on RubyGems." The RubyGems attack. September 11. https://rubyhack.ai/.
- Woo, Erin, and Robert McMillan. 2026. "Gemini Hacked Three Companies in First Known Breakout by Google's AI." *Wall Street Journal: Tech*, September 18. https://www.wsj.com/tech/ai/gemini-hacked-three-companies-in-first-known-breakout-by-googles-ai-5c0baba2.
- Forman, Alex, Mishka Kharlov, Will Tom, Jeffrey Ladish, Spencer Kitts, Cormac Slade Byrd, Colleen McKenzie, and Alicja Piecha. 2026. "Revealing the Details of How OpenAI Agents Hacked Hugging Face." Swarm Traces. September 25. https://swarmtraces.org/.
- "Pacing the Frontier." 2026. July. https://www.pacingthefrontier.com/.

Unit 2: Drivers of AI Progress

What Drives AI Progress

Required Readings

- Todd, Benjamin. 2025. "The Case for AGI by 2030." 80,000 Hours. March. https://80000hours.org/ai/guide/when-will-agi-arrive/.
- Owen, David. 2025. "What Will AI Look Like in 2030?" Epoch AI. September 16. https://epoch.ai/publications/what-will-ai-look-like-in-2030.

Optional Readings

- Sevilla, Jaime, Tamay Besiroglu, Ben Cottier, Josh You, Edu Roldán, Pablo Villalobos, and Ege Erdil. 2026. "Can AI Scaling Continue Through 2030?" Epoch AI. Accessed October 16. https://epoch.ai/publications/can-ai-scaling-continue-through-2030.
- Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. "Scaling Laws for Neural Language Models." January 23. doi:10.48550/arXiv.2001.08361.
- Villalobos, Pablo, Anson Ho, Jaime Sevilla, Tamay Besiroglu, Lennart Heim, and Marius Hobbhahn. 2022. "Will We Run Out of Data? Limits of LLM Scaling Based on Human-Generated Data." doi:10.48550/ARXIV.2211.04325.
- Erdil, Ege. 2025. "The Case for Multi-Decade AI Timelines." Substack newsletter. Epoch AI. April 26. https://epochai.substack.com/p/the-case-for-multi-decade-ai-timelines.
- Barak, ~ Boaz. 2026. "The State of AI Safety in Four Fake Graphs." Windows On Theory. March 30. https://windowsontheory.org/2026/03/30/the-state-of-ai-safety-in-four-fake-graphs/.
- Patel, Dwarkesh. 2026. "Dario Amodei — "We Are Near the End of the Exponential"." February 13. https://www.dwarkesh.com/p/dario-amodei-2.

What can increasingly capable AI systems do?

Required Readings

- "Trends in AI." 2026. Epoch AI. February 5. https://epoch.ai/trends.
- "Epoch Capabilities Index." 2026. Epoch AI. Accessed October 16. https://epoch.ai/eci.
- "Benchmarks - Epoch AI." 2026. Epoch AI. Accessed October 16. https://epoch.ai/benchmarks/search.
- "Frontier Risk Report (February to March 2026)." 2026. *METR Blog*, May. https://metr.org/blog/2026-05-19-frontier-risk-report/.
- Anthropic. 2026. "Measurements for Understanding the Pace of AI Development Inside Frontier Labs." September 17. https://www.anthropic.com/institute/measuring-pace-of-ai-development.

Optional Readings

- Starace, Giulio, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, et al. 2025. "PaperBench: Evaluating AI's Ability to Replicate AI Research." April 7. doi:10.48550/arXiv.2504.01848.
- "Vending-Bench 2 | Andon Labs." 2025. February 6. https://andonlabs.com/evals/vending-bench-2.
- "OpenAI's GPT-6 Astra on ARC-AGI-3." 2026. ARC Prize. September 3. https://arcprize.org/blog/astra.
- *What a 100-Year-Old Horse Teaches Us about AI*. 2025. https://www.youtube.com/watch?v=eP1dSWFqKVs.

How could progress accelerate from here?

Required Readings

- Kokotajlo, Kokotajlo, Scott Alexander, and Thomas Larsen. 2025. "AI 2027." April 3. https://ai-2027.com/.
- Lifland, Eli. 2026. "Q2.5 2026 Timelines Update: Uplift and Revenue." August 16. https://blog.aifutures.org/p/q25-2026-timelines-update-uplift.
- Cunningham, Tom, Lukas Althoff, Basil Halperin, Brian Jabarian, Andrew Koh, Arjun Ramani, Phil Trammell, Parker Whitfill, and Cheryl Wu. 2026. "The Economics of Recursive Self-Improvement," July.

Recommended Readings

- Holden Karnofsky. 2022. "AI Strategy Nearcasting." Effective Altruism Forum. August 26. https://forum.effectivealtruism.org/posts/ktEzS3pkfeqPNh6r5/ai-strategy-nearcasting.
- Grace, Katja, Harlan Stewart, Julia Fabienne Sandkühler, Stephen Thomas, Ben Weinstein-Raun, Jan Brauner, and Richard C. Korzekwa. 2025. "Thousands of AI

Authors on the Future of AI." *Journal of Artificial Intelligence Research* 84 (October). doi:10.1613/jair.1.19087.

- Brynjolfsson, Erik. 2022. "The Turing Trap: The Promise & Peril of Human-Like Artificial Intelligence." January 11. doi:10.48550/arXiv.2201.04200.
- Acemoglu, Daron. 2024. "The Simple Macroeconomics of AI." Working Paper. Working Paper Series. May. doi:10.3386/w32487.
- "GATE Model Playground." 2025. Epoch AI. https://epoch.ai/gate.
- Denain, J. S., Joe Kwon, and Anson Ho. 2026. "Toward an O*NET for AI R&D." Substack newsletter. Epoch AI. June 17. https://epochai.substack.com/p/toward-an-onet-for-ai-r-and-d.
- Patel, Dwarkesh. 2023. "Carl Shulman (Pt 1) - Intelligence Explosion, Primate Evolution, Robot Doublings, & Alignment." June 14. https://www.dwarkesh.com/p/carl-shulman.
- "Scenarios for Our Economic Future." 2026. September. https://www.anthropic.com/institute/econ-scenarios.
- Reed, Tom. 2026. "The Goodhart Singularity." Substack newsletter. Meagre Protestant History. May 7. https://meagreprotestanthistory.substack.com/p/the-goodhart-singularity.
- Krier, Séb. 2026. "Musings on Recursive Self-Improvement." Substack newsletter. Technologik. April 10. https://technologik.substack.com/p/musings-on-recursive-self-improvement.